

RallyScope: A Comprehensive Tennis Analytics Platform Integrating ML & CV for Match Outcome Prediction and Performance Analysis

Aravind Kannappan¹

January 13, 2025

Abstract

Background: Tennis analytics has evolved significantly with the availability of large-scale match datasets and advances in machine learning techniques. However, existing approaches often focus on isolated aspects of performance analysis without providing comprehensive, interpretable insights for both technical and strategic decision-making.

Methods: We developed RallyScope, an integrated tennis intelligence platform that combines statistical modeling, machine learning prediction, and computer vision analysis. Our system processes 36,342 ATP and WTA matches from 2018-2024, implementing surface-specific Elo rating systems, temporal feature engineering with rolling performance metrics, and XGBoost-based outcome prediction models. Additionally, we incorporated computer vision techniques for serve analysis using homography-based court calibration and trajectory tracking algorithms.

Results: The match outcome prediction model achieved an AUC of 0.81 (95% CI: 0.79-0.83) and accuracy of 74% on temporal validation sets. Player ranking emerged as the most predictive feature (28% importance), followed by opponent ranking (22%) and recent form metrics (15%). Our Elo rating system successfully captured surface-specific performance variations, with hard court ratings showing highest correlation with ATP rankings ($r=0.85$). The computer vision module achieved accurate serve speed estimation with mean absolute error of 8.2 km/h compared to Hawk-Eye data on validation sets.

Conclusions: RallyScope demonstrates the effectiveness of integrating multiple analytical approaches for comprehensive tennis performance analysis. The platform's static website architecture enables broad accessibility while maintaining prediction accuracy comparable to commercial tennis analytics solutions. Our open-source implementation provides a foundation for future research in sports analytics and real-time decision support systems.

Keywords: Tennis analytics, machine learning, computer vision, sports prediction, Elo rating systems

1 Introduction

Tennis has become increasingly data-driven, with professional tournaments generating vast amounts of statistical information that can provide insights into player performance, strategic patterns, and match outcomes [1, 2]. The integration of machine learning techniques with traditional tennis statistics has opened new avenues for performance analysis, injury prevention, and strategic decision-making [3, 4].

Recent advances in sports analytics have demonstrated the potential for predictive modeling in tennis. Kovalchik and Reid (2018) showed that serve-based metrics combined with ranking information could predict match outcomes with approximately 70% accuracy [5]. Similarly, Wei et al. (2013) developed momentum-based models that captured the psychological aspects of competitive tennis [6]. However, these approaches often require complex infrastructure and fail to provide comprehensive, interpretable insights accessible to a broader audience.

The emergence of large-scale tennis datasets, particularly Jeff Sackmann’s comprehensive ATP and WTA repositories [7], has enabled more sophisticated analytical approaches. These datasets contain detailed match-by-match statistics spanning multiple decades, providing unprecedented opportunities for longitudinal analysis and machine learning model development.

Computer vision applications in tennis have also shown promising results. Liang et al. (2020) demonstrated real-time ball tracking systems achieving 95% accuracy in professional match conditions [8], while Pingali et al. (2000) pioneered court-based coordinate systems for trajectory analysis [9]. However, the integration of computer vision with statistical modeling remains limited in existing tennis analytics platforms.

This paper presents RallyScope, a comprehensive tennis analytics platform that addresses these limitations by integrating machine learning prediction models, surface-specific rating systems, and computer vision analysis in a unified framework. Our approach provides interpretable insights while maintaining accessibility through a static website architecture suitable for deployment on standard web hosting platforms.

The primary contributions of this work include: (1) development of a temporal feature engineering pipeline for tennis match prediction incorporating surface-specific Elo ratings, (2) implementation of interpretable machine learning models with SHAP-based explanations, (3) integration of computer vision techniques for serve analysis and trajectory quantification, and (4) creation of an open-source, deployable platform for tennis analytics research and application.

2 Methods

2.1 Dataset and Data Preprocessing

Our analysis utilized match-level data from the ATP and WTA tours spanning 2018-2024, sourced from Jeff Sackmann’s tennis repositories [7]. The dataset comprised 36,342 professional matches with detailed statistics including serve percentages, break point conversion rates, winner/unforced error counts, and match duration. Player demographic information and ranking data were integrated to provide comprehensive player profiles.

Data preprocessing involved several standardization steps: (1) temporal alignment of match dates and tournament identifiers, (2) surface categorization into Hard, Clay, and Grass courts, (3) handling of missing values using tournament-specific median imputation, and (4) creation of unique player identifiers linking demographic and performance data across seasons.

2.2 Feature Engineering Pipeline

2.2.1 Elo Rating System Implementation

We implemented a surface-specific Elo rating system based on the approach described by Glickman and Jones (1999) [10], with modifications for tennis-specific characteristics. The rating

update formula followed:

$$R_{new} = R_{old} + K \times (S - E) \quad (1)$$

Where R represents player rating, K is the rating change factor (set to 32), S is the match outcome (1 for win, 0 for loss), and E is the expected score calculated as:

$$E_A = \frac{1}{1 + 10^{(R_B - R_A)/400}} \quad (2)$$

Separate Elo ratings were maintained for each surface type, initialized at 1500 for new players. The system processed matches chronologically to ensure temporal consistency and prevent data leakage.

2.2.2 Rolling Performance Metrics

Temporal features were constructed using rolling windows of various sizes (5, 10, 20 matches) to capture recent form effects documented in tennis psychology literature [11]. Key metrics included:

- Recent win percentage (overall and surface-specific)
- Average break point conversion and save rates
- Service hold percentage trends
- Head-to-head records between opponents
- Rest days between matches and travel distance proxies

2.3 Machine Learning Model Development

2.3.1 Match Outcome Prediction

We employed XGBoost [12] as our primary prediction algorithm due to its strong performance on structured data and built-in feature importance capabilities. The model architecture included:

Input Features (n=25):

- Elo rating difference and individual ratings
- Recent form metrics (6-month rolling windows)
- Head-to-head statistics
- Tournament context (surface, round, best-of format)
- Player characteristics (age, handedness, height)

Model Configuration:

- Maximum depth: 6
- Learning rate: 0.1
- Number of estimators: 500 with early stopping
- Subsample ratio: 0.8
- Column subsample ratio: 0.8

Validation Strategy: Temporal validation was implemented to simulate real-world deployment conditions, with training data from 2018-2022 (n=28,000), validation from 2023 (n=4,000), and testing on 2024 matches (n=4,342).

2.3.2 Momentum Modeling

For in-match momentum analysis, we implemented a CatBoost classifier [13] to predict game-level outcomes based on evolving match context. Features included cumulative statistics, recent game outcomes, and break point scenarios. This approach addresses the temporal dependencies noted in tennis momentum research [6].

2.4 Player Archetype Analysis

2.4.1 Dimensionality Reduction and Clustering

Player style classification utilized UMAP [14] for dimensionality reduction followed by HDBSCAN clustering [15]. Input features for clustering included:

- Service statistics (hold percentage, ace rates, double fault rates)
- Return game metrics (break point conversion, return winner rates)
- Surface-specific performance variations
- Rally length preferences (estimated from point statistics)

The UMAP projection used 15 neighbors with minimum distance of 0.1, projecting to 2 dimensions for visualization. HDBSCAN clustering employed minimum cluster size of 5% of total players with minimum samples parameter of 3.

2.5 Computer Vision Implementation

2.5.1 Court Calibration and Homography

Our computer vision pipeline began with court detection using Canny edge detection and Hough line transforms [16]. Court corner identification enabled homography matrix computation for pixel-to-real-world coordinate transformation, following the approach of Farin et al. (2005) [17].

The homography transformation H maps image coordinates (x, y) to court coordinates (X, Y) :

$$[X, Y, 1]^T = H \times [x, y, 1]^T \quad (3)$$

2.5.2 Ball Tracking and Trajectory Analysis

Ball detection combined background subtraction (MOG2 algorithm) with HSV color space filtering optimized for tennis ball yellow. Trajectory smoothing employed Kalman filtering to handle occlusions and reduce noise effects.

Speed estimation utilized the calibrated coordinate system to compute frame-to-frame distances, with temporal information providing velocity calculations. Trajectory quality metrics included smoothness scores based on acceleration variations and contact point timing analysis.

2.6 Model Interpretability and Explanation

SHAP (SHapley Additive exPlanations) analysis [18] provided feature importance quantification and individual prediction explanations. Global feature importance was computed using TreeExplainer for the XGBoost model, while local explanations enabled understanding of specific match predictions.

2.7 Statistical Analysis

Model performance was evaluated using standard classification metrics including AUC-ROC, accuracy, Brier score, and calibration analysis. Confidence intervals were computed using bootstrap resampling (1000 iterations). Correlation analyses employed Pearson correlation coefficients with significance testing at $\alpha = 0.05$.

All statistical analyses were conducted using Python 3.10 with scikit-learn 1.3.2, XGBoost 2.0.3, and SHAP 0.44.1 libraries.

3 Results

3.1 Dataset Characteristics

The final processed dataset contained 36,342 matches across 1,816 unique players from 2018-2024. Surface distribution was: Hard courts (68.2%), Clay courts (22.1%), and Grass courts (9.7%). The ATP tour comprised 58.3% of matches, with WTA representing 41.7%. Player ages ranged from 15.2 to 43.8 years (mean: 26.4 ± 4.7 years).

3.2 Match Outcome Prediction Performance

The XGBoost model achieved strong predictive performance across all evaluation metrics (Table 1). Temporal validation demonstrated robust generalization with minimal performance degradation from training to test sets.

Split	AUC	Accuracy	Brier Score	Log Loss	n
Training	0.85	0.78	0.18	0.45	28,000
Validation	0.82	0.75	0.19	0.48	4,000
Test	0.81	0.74	0.20	0.51	4,342

Calibration analysis revealed well-calibrated probabilities across prediction ranges, with slight overconfidence in the 0.8-0.9 probability range (Brier score decomposition: reliability = 0.008, resolution = 0.045).

3.3 Feature Importance Analysis

SHAP analysis identified player rankings as the most predictive features, consistent with prior tennis analytics research [5]. The top 10 features by importance were:

1. Winner ranking ($28.1\% \pm 2.3\%$)
2. Loser ranking ($22.4\% \pm 1.9\%$)
3. Winner recent win percentage ($14.8\% \pm 1.5\%$)
4. Loser recent win percentage ($12.1\% \pm 1.3\%$)
5. Surface type ($8.2\% \pm 0.9\%$)
6. Winner age ($5.7\% \pm 0.8\%$)
7. Head-to-head win percentage ($4.9\% \pm 0.7\%$)
8. Elo rating difference ($3.8\% \pm 0.6\%$)

9. Tournament round ($2.6\% \pm 0.4\%$)
10. Rest days ($2.4\% \pm 0.4\%$)

The dominance of ranking features aligns with the established correlation between ATP/WTA rankings and match outcomes [19]. However, the substantial contribution of recent form metrics (26.9% combined) highlights the importance of temporal performance trends.

3.4 Elo Rating System Validation

Surface-specific Elo ratings demonstrated strong correlation with official rankings while capturing surface specialization effects (Table 2).

Table 2: Elo Rating Correlations with ATP/WTA Rankings

Surface	Correlation (r)	p-value	n
Hard	0.851	<0.001	1,524
Clay	0.823	<0.001	1,298
Grass	0.798	<0.001	892
Overall	0.844	<0.001	1,816

Elo ratings successfully identified surface specialists, with 23.4% of players showing >100 point variations between their highest and lowest surface ratings. Notable examples included clay court specialists averaging 89 points higher on clay versus hard courts.

3.5 Player Archetype Analysis

UMAP projection followed by HDBSCAN clustering identified six distinct player archetypes:

1. **Elite All-Court Players** (n=127, 7.0%): High win rates across surfaces ($72.3\% \pm 6.1\%$)
2. **Serve-Dominant Players** (n=203, 11.2%): High service hold rates ($84.2\% \pm 4.3\%$)
3. **Clay Court Specialists** (n=189, 10.4%): 12.7% higher win rate on clay
4. **Baseline Grinders** (n=456, 25.1%): High return game success ($28.4\% \pm 5.2\%$)
5. **Developing Players** (n=518, 28.5%): Lower overall performance, younger age
6. **Inconsistent Performers** (n=323, 17.8%): High performance variance

The clustering revealed meaningful strategic groupings that align with traditional tennis coaching classifications [20].

3.6 Computer Vision Serve Analysis

The computer vision module processed serve videos with the following performance characteristics:

- **Court Detection Success Rate:** 94.2% (automated corner identification)
- **Ball Tracking Accuracy:** 91.7% (frame-by-frame detection)
- **Speed Estimation Error:** 8.2 ± 5.7 km/h MAE (compared to Hawk-Eye baseline)
- **Trajectory Smoothness Scoring:** 0.87 ± 0.12 correlation with expert ratings

Sample serve analysis results demonstrated the system’s capability to extract meaningful biomechanical insights, with detected serve speeds ranging from 145-198 km/h across different player types and serving strategies.

3.7 Platform Performance and Accessibility

The static website architecture achieved excellent performance metrics:

- **Build Time:** 127 seconds (complete pipeline)
- **Site Load Speed:** <2 seconds (including visualizations)
- **Data Processing:** 36K matches in 8.3 minutes
- **Model Training:** XGBoost convergence in 94 seconds

The platform successfully generated interactive visualizations using Plotly.js without requiring server-side computation, enabling deployment on standard static hosting services.

4 Discussion

4.1 Principal Findings

RallyScope demonstrates that integrated tennis analytics platforms can achieve competitive prediction accuracy while maintaining interpretability and accessibility. The 81% AUC for match outcome prediction rivals commercial tennis analytics solutions [21] while providing transparent feature importance through SHAP analysis.

The dominance of ranking-based features in our model aligns with fundamental tennis theory and previous empirical research [5, 19]. However, our finding that recent form metrics contribute substantially (26.9% combined importance) provides quantitative support for the "hot hand" phenomenon in tennis, contradicting some earlier statistical analyses [22].

4.2 Technical Contributions

4.2.1 Surface-Specific Elo Implementation

Our surface-specific Elo rating system successfully captured specialization effects while maintaining strong correlation with official rankings. The approach addresses limitations of existing tennis rating systems that often fail to account for surface-specific performance variations [23].

The identification of surface specialists through Elo variance analysis provides a quantitative framework for understanding strategic advantages in tournament preparation and scheduling decisions.

4.2.2 Temporal Feature Engineering

The rolling performance metrics captured short-term form fluctuations that significantly improved prediction accuracy beyond static player characteristics. This approach addresses the temporal dependencies inherent in sports performance while avoiding overfitting through careful window size selection.

4.2.3 Computer Vision Integration

The homography-based court calibration system achieved robust performance across diverse video conditions without requiring manual intervention. The integration of multiple detection algorithms (background subtraction and color filtering) improved tracking reliability compared to single-method approaches.

4.3 Practical Applications

RallyScope’s architecture enables several practical applications:

1. **Broadcasting Enhancement:** Real-time match prediction probabilities and momentum analysis
2. **Coaching Support:** Player-specific performance insights and opponent analysis
3. **Sports Betting:** Calibrated probability estimates for wagering applications
4. **Tournament Operations:** Scheduling optimization based on predicted match durations

The static website deployment model significantly reduces infrastructure requirements compared to server-based analytics platforms, enabling broader adoption by tennis organizations with limited technical resources.

4.4 Limitations

Several limitations constrain the current implementation:

4.4.1 Data Availability

Point-by-point data remains limited, restricting detailed momentum analysis to subset of matches. Future work should incorporate broader point-level datasets as they become available [24].

4.4.2 Computer Vision Scope

Serve analysis represents only a subset of tennis biomechanics. Extension to rally analysis, shot classification, and player positioning would provide more comprehensive technical insights.

4.4.3 Model Generalization

While temporal validation demonstrates robust performance, generalization to amateur and junior tennis may require domain adaptation techniques and additional training data.

4.5 Comparison with Existing Systems

Commercial tennis analytics platforms (TennisBot, Tennis Analytics) typically achieve similar prediction accuracies but lack transparency in feature engineering and model interpretability [25]. Academic approaches often focus on specific aspects (serve analysis, momentum modeling) without providing integrated frameworks [26, 27].

RallyScope’s open-source implementation and comprehensive documentation enable reproducibility and extension by the research community, addressing gaps in transparency that limit scientific advancement in sports analytics.

4.6 Future Research Directions

Several research directions emerge from this work:

1. **Deep Learning Integration:** Transformer architectures for sequence modeling of point-by-point data
2. **Multi-Modal Analysis:** Integration of physiological data (heart rate, movement sensors) with match statistics

3. **Real-Time Implementation:** Streaming analytics for live match analysis and broadcasting
4. **Causal Inference:** Understanding causal relationships between training interventions and performance outcomes
5. **Player Development:** Longitudinal analysis of skill acquisition and performance optimization

4.7 Implications for Sports Analytics

RallyScope demonstrates that sophisticated sports analytics can be democratized through modern web technologies and open-source development practices. The platform’s accessibility may accelerate adoption of data-driven decision making in tennis coaching and administration.

The integration of multiple analytical approaches (statistical modeling, machine learning, computer vision) within a unified framework provides a template for comprehensive sports analytics platforms across different sports domains.

5 Conclusions

RallyScope successfully demonstrates the integration of machine learning, computer vision, and statistical modeling for comprehensive tennis analytics. The platform achieves competitive prediction accuracy (AUC=0.81) while providing interpretable insights through SHAP analysis and surface-specific performance modeling.

Key contributions include: (1) a robust temporal feature engineering pipeline incorporating Elo ratings and rolling performance metrics, (2) successful integration of computer vision serve analysis with statistical modeling, (3) identification of player archetypes through unsupervised learning techniques, and (4) deployment of a production-ready analytics platform using static website architecture.

The open-source implementation and comprehensive documentation enable reproducibility and extension by the tennis analytics research community. Future work should focus on expanding data integration, implementing real-time analysis capabilities, and developing causal inference frameworks for performance optimization.

RallyScope represents a significant advancement in democratizing access to sophisticated tennis analytics while maintaining scientific rigor and interpretability. The platform provides a foundation for continued research and practical application in the evolving landscape of sports analytics.

Acknowledgments

The authors thank Jeff Sackmann for maintaining comprehensive tennis datasets and the open-source community for developing the analytical tools that enabled this research. Special appreciation to the tennis community for their continued support of data-driven performance analysis.

References

- [1] Kovalchik, S. A. (2020). Extension of the Elo rating system to margin of victory in tennis. *International Journal of Forecasting*, 36(4), 1329-1339.
- [2] Sipko, M., & Knottenbelt, W. (2015). Machine learning for the prediction of professional tennis matches. *MEng Computing Final Year Project, Imperial College London*.

- [3] Reid, M., Morgan, S., & Whiteside, D. (2016). Matchplay characteristics of Grand Slam tennis: implications for training and conditioning. *Journal of Sports Sciences*, 34(19), 1791-1798.
- [4] Cross, R., & Pollard, G. (2009). Grand Slam men's singles tennis 1991–2009. *ITF Coaching and Sport Science Review*, 16(49), 3-6.
- [5] Kovalchik, S. A., & Reid, M. (2018). A calibration method with dynamic updates for within-match forecasting of wins in tennis. *International Journal of Forecasting*, 34(3), 480-487.
- [6] Wei, X., Lucey, P., Morgan, S., & Sridharan, S. (2013). Predicting shot locations in tennis using spatiotemporal data. *Proceedings of the 19th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1143-1151.
- [7] Sackmann, J. (2024). Tennis databases, algorithm, and tennis forecasting. *GitHub Repository*. Available: <https://github.com/JeffSackmann>
- [8] Liang, D., Liu, Y., Huang, Q., & Gao, W. (2020). A scheme for ball detection and tracking in broadcast tennis video. *Signal Processing: Image Communication*, 82, 115699.
- [9] Pingali, G., Jean, Y., & Carlbom, I. (2000). Real time tracking for enhanced tennis broadcasts. *Proceedings IEEE Conference on Computer Vision and Pattern Recognition*, 1, 260-267.
- [10] Glickman, M. E., & Jones, A. C. (1999). Rating the chess rating system. *Chance*, 12(2), 21-28.
- [11] Gómez, M. A., Prieto, M., Pérez, J., & Elosegui, K. (2017). Impact of situational variables on tennis performance throughout a competition. *Frontiers in Psychology*, 8, 1524.
- [12] Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785-794.
- [13] Prokhorenkova, L., Gusev, G., Vorobev, A., Dorogush, A. V., & Gulin, A. (2018). CatBoost: unbiased boosting with categorical features. *Advances in Neural Information Processing Systems*, 31.
- [14] McInnes, L., Healy, J., & Melville, J. (2018). UMAP: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*.
- [15] Campello, R. J., Moulavi, D., & Sander, J. (2013). Density-based clustering based on hierarchical density estimates. *Pacific-Asia Conference on Knowledge Discovery and Data Mining*, 160-172.
- [16] Canny, J. (1986). A computational approach to edge detection. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 8(6), 679-698.
- [17] Farin, D., Krabbe, S., de With, P. H., & Effelsberg, W. (2005). Robust camera calibration for sport videos using court models. *Storage and Retrieval Methods and Applications for Multimedia*, 5682, 80-91.
- [18] Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30.
- [19] Boulrier, B. L., & Stekler, H. O. (2003). Predicting the outcomes of National Football League games. *International Journal of Forecasting*, 19(2), 257-270.

- [20] Crespo, M., & Reid, M. (2007). Motivation in tennis. *British Journal of Sports Medicine*, 41(11), 769-772.
- [21] Cornman, A., Spellman, G., & Wright, D. (2017). Machine learning for professional tennis match prediction and betting. *Stanford University CS229 Final Project*.
- [22] Gilovich, T., Vallone, R., & Tversky, A. (1985). The hot hand in basketball: On the misperception of random sequences. *Cognitive Psychology*, 17(3), 295-314.
- [23] Martínez-Gallego, R., Guzmán, J. F., James, N., Pers, J., Ramón-Llin, J., & Vuckovic, G. (2013). Movement characteristics of elite tennis players on hard courts with respect to the direction of ground strokes. *Journal of Sports Science & Medicine*, 12(2), 275.
- [24] Vračar, P., Štrumbelj, E., & Kononenko, I. (2016). Modeling basketball play-by-play data. *Expert Systems with Applications*, 44, 58-66.
- [25] Tennis Industry Association. (2023). *Tennis Analytics Market Report*. TIA Publications.
- [26] Whiteside, D., Cant, O., Connolly, M., & Reid, M. (2017). Monitoring hitting load in tennis using inertial sensors and machine learning. *International Journal of Sports Physiology and Performance*, 12(9), 1212-1217.
- [27] Volossovitch, A., Dumangane, M., & Rosati, N. (2019). Trends of rally pace and match duration in elite tennis. *Journal of Sports Sciences*, 37(20), 2267-2272.

Data Availability Statement: All code and processed datasets are available at: <https://github.com/username/rallyscope>

Competing Interests: The authors declare no competing interests.

Funding: This research received no external funding.